Decision Stumps

- A decision stump is a basic form of a decision tree, consisting of a single root node connected to two leaf nodes. It splits the dataset on one feature.

Example:

- Root Node: Splits data based on a numerical threshold.
- Left Child Node: Contains instances where the feature value is <= 50.
- Right Child Node: Contains instances where the feature value is > 50.

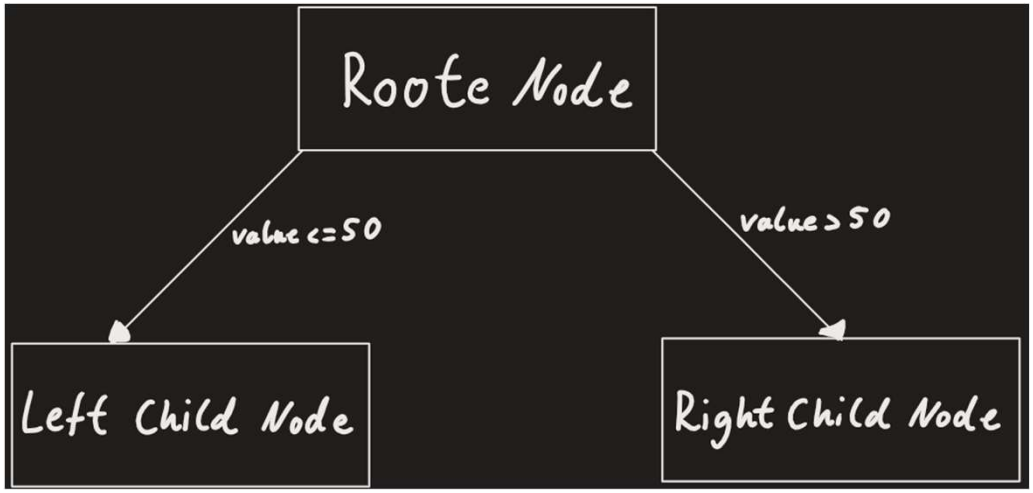

Entropy and Information Gain

Entropy:

- Entropy measures the impurity or disorder within a dataset. It helps to determine how a feature divides the data into groups that are homogeneous with respect to the target attribute.
- Formula: $H = -\sum_i p(c_i)\ log_2\ p(c_i)$, where $p(c_i)$ is the probability of class c_i.
- A high entropy value indicates a mixture of different classes, whereas a low entropy suggests a homogenous set.

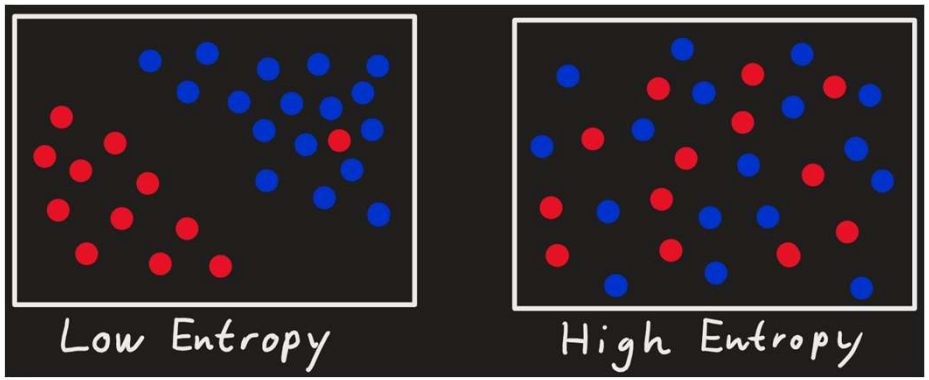

Information Gain:

- Measures the reduction in entropy after a dataset is split on an attribute. It is crucial for determining the best attribute that yields the most informative split at each node in a tree.
- Formula: $\Delta H = H_p\ - [\frac{n_1}{n} H_1 + \frac{n_2}{n} H_2]$, where H_p is the entropy of the parent set before split, and H_1, H_2 are the entropies of the two sets after the split.
- A higher Information Gain value signifies a more effective attribute for splitting the data, leading to purer nodes.

Decision Stump Implementation

1. Data Preparation

- Load data into DataFrame.
- Initialize an empty list for results.

2. Feature Analysis

- Iterate over each feature, excluding the target ('class').
- Sort values and split data to calculate information gain and entropy.

3. Result Compilation

- Store results (feature, value, information gain, entropy) in a DataFrame and print it.